

Walchand College of Engineering, Sangli
Department of CSE

Seminar on
“Scaling Up Machine Learning and Deep
Learning : Parallel Approach with CUDA
and OpenMP”

Roll No.: 2018MTECSCO007

Name : Veersen Vijay Jadhav

Contents


Sr. No.	Title	Slide No.
01	Research Area	03
02	Technology	04
03	Why this topic ?	05
04	Literature Survey	06
05	Methodology	09
06	Implementation	17
07	Result & Analysis	20
08	Conclusion	24
09	References	25

Research Area

- ▶ Research Area : Applying High Performance Computing / Parallel Computing to Machine Learning Algorithms.
- ▶ Paper Title : Accelerating nearest neighbour partitioning neural network classifier based on CUDA.
- ▶ Authors : Lin Wang, Ajith Abraham, Meihui Li
- ▶ Publisher : Elsevier
- ▶ Journal : Engineering Applications of Artificial Intelligence
- ▶ Publication Year : 2018

Technology

WORKING DISTRIBUTION



14 April 2019

Why Machine Learning with High Performance Computing ?

- ▶ Numerical analysis formed backbone for supercomputing devices over the decades.
- ▶ Recently scientists have begun experimenting with understanding complex systems using
 - ▶ Machine Learning predictive models
 - ▶ Deep Neural Network
 - ▶ With Parallel ComputingTrained by virtually unlimited datasets produced globally.
- ▶ HPC can improve accuracy, accelerate time to solution and significantly reduce costs.

Literature Survey

- ▶ Anderson, D., Coupland, S., **Parallelisation of fuzzy inference on a graphics processor unit using the compute unified device architecture.** Recent. Progr. Med. 85 (3), 160-165, 2008.
- ▶ Jang, H., Park, A., Jung, K., **Neural network implementation using cuda and openmp.** In: **Digital Image Computing: Techniques and Applications.** pp. 155-161, 2008.
- ▶ Wang, L., Yang, B., Chen, Y., **Improving particle swarm optimization using multilayer searching strategy.** Inform. Sci. 274 (8), 70-94, 2014.
- ▶ Wang, L., Yang, B., Chen, Y., Abraham, A., Sun, H., Chen, Z., Wang, H., **Improvement of neural network classifier using floating centroids.** Knowl. Inf. Syst. 31 (3), 433-454, 2012.

Literature Survey : Floating Centroid Method

Fixed Centroids

- In traditional Neural Network, features of centroids are fixed.
- In mapping of samples, fixed centroids reduces possibilities of finding optimal solution.

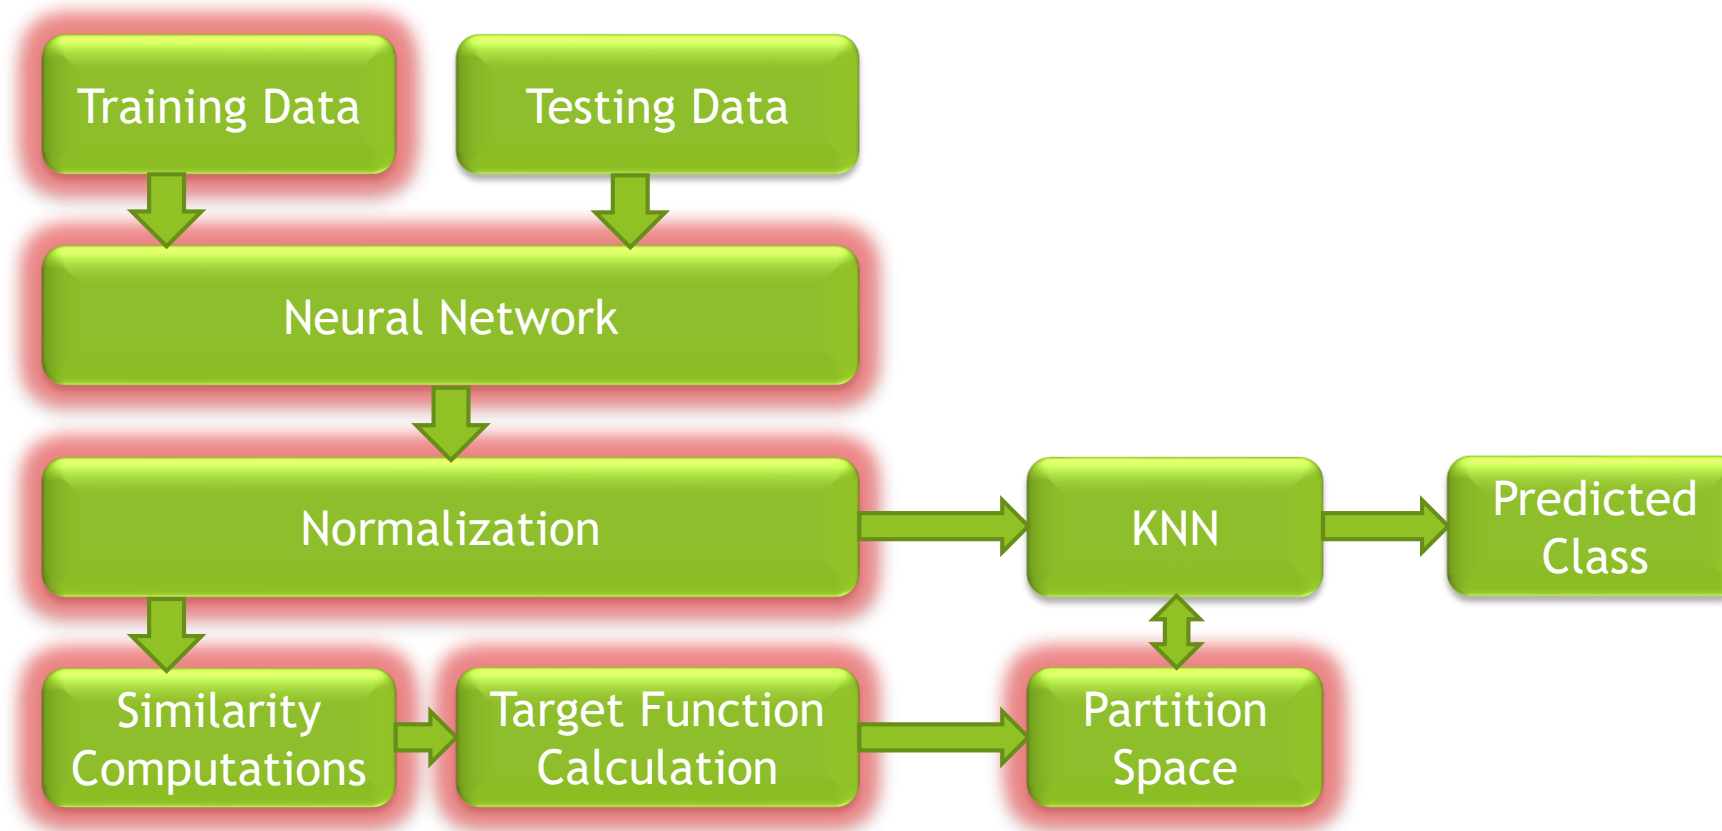
Floating Centroid Method

- In partition space, centroid is point. The proposed approach introduces many floating centroids.
- FCMs are spread throughout partition space, and obtained by using K-Means algorithms.

Limitations of FCM

- Cannot yield flexible decision boundaries.
- Problem overcomes by Neural Network Partitioning.

Literature Survey : Particle Swarm Optimization



Methodology

Sample Mapping

Normalization

Similarity Computations

Target Function Calculation

Sample Mapping in Parallel

In CUDA, data transmission takes place between CPU and GPU

Original samples → Global Memory
Information of Neural Network → Shared Memory

Data Divided as →

$$|S^i| = \begin{cases} \left\lceil \frac{|S|}{threadcount} \right\rceil, & i < threadcount \\ |S| \bmod \left\lceil \frac{|S|}{threadcount} \right\rceil, & i = threadcount \end{cases}$$

Sample Mapping in Parallel

$tx \rightarrow$ identity of thread block

Value of identity of particle \rightarrow

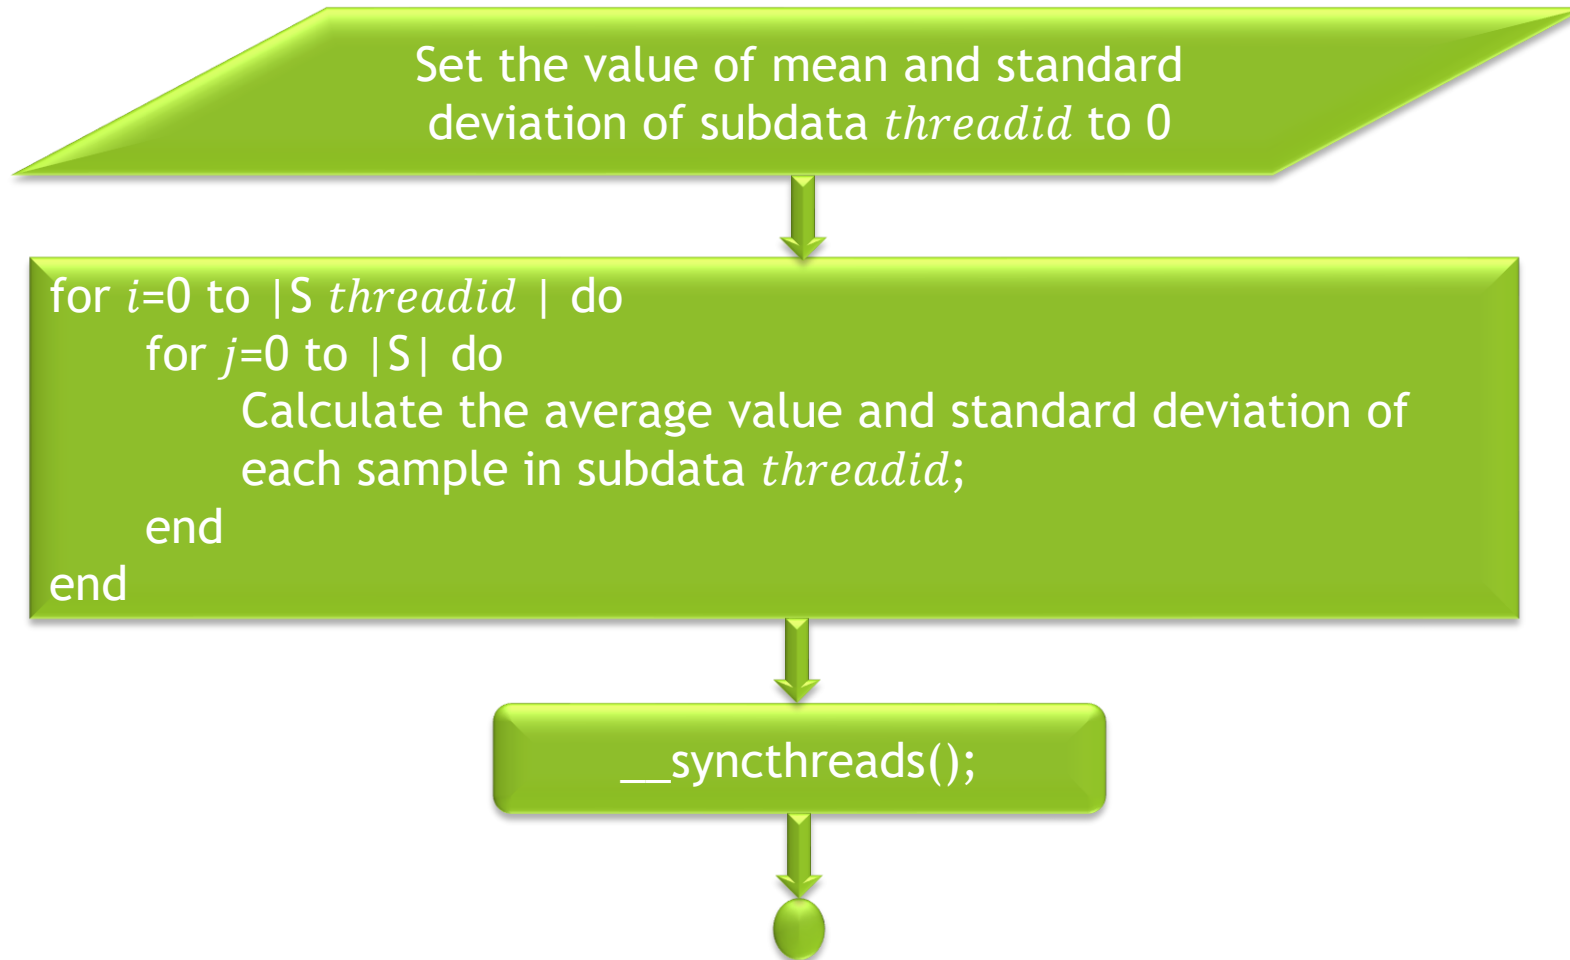
$$particleid = tx / threadcount$$

The value of identity of thread in each neural network \rightarrow

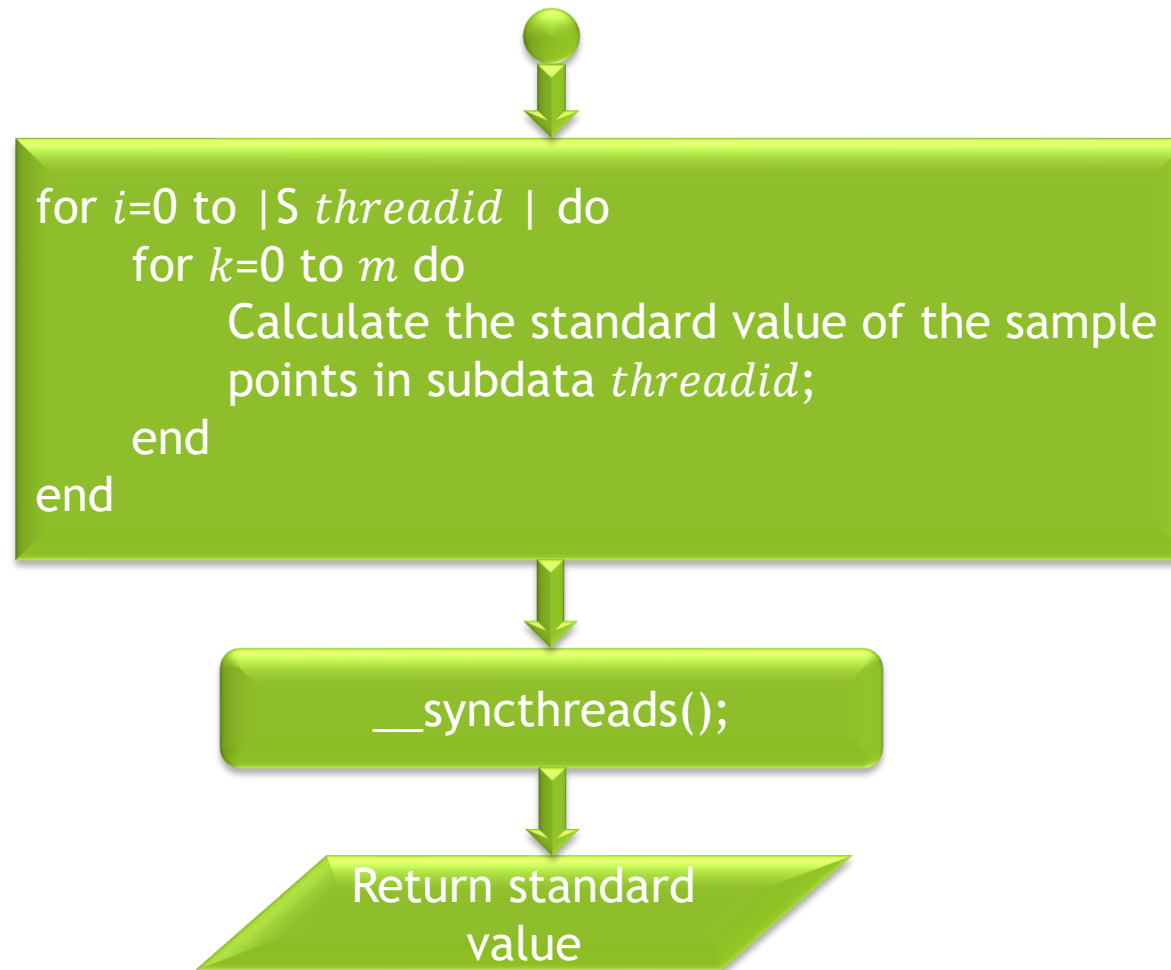
$$threadid = tx \bmod threadcount$$

Evaluation of sample is executed parallelly by tasks $\{T^1, T^2, \dots, T^{threadnum}\}$ are allocated to the $threadnum$ parallel thread in one block. Then mapped samples stored in global memory.

Normalization in Parallel



Normalization in Parallel




Similarity Computation in Parallel

$$d = |x| \frac{1 - e^{-|x|/2}}{|x| + |x| \cdot e^{-|x|/2}}$$

```
for  $i=0$  to  $|S_{threadid}|$  do  
  for  $j=0$  to  $m$  do  
    Calculate  $d$  use formula in subdata  
     $threadid$ ; Obtain the values by the  
    normalized points into a hypersphere;  
  end  
end
```

__syncthreads();

Similarity Computation in Parallel



```
for  $i=0$  to  $|S_{threadid}|$  do
  for  $j=0$  to  $|S|$  do
    Calculate the distance  $D$  between two points with
    Euclidean Distance in subdata  $threadid$ ;
    if  $threadid=j$  then
      The value of  $SimilarityArray_{threadid,j}$  is 2;
    else
      The value of  $SimilarityArray_{threadid,j}$  is  $2-D$ ;
    end
  end
end
```



```
__syncthreads();
```



```
Return standard  
value
```

Target Function Calculation in Parallel

$$F = \omega(x_i)(S_{nonself}(x_i) - \alpha S_{self}(x_i))$$

$\alpha \rightarrow$ adjustment coefficient

$S_{nonself}(x_i) \rightarrow$ Sum of similarities between x_i and samples in other classes

$S_{self}(x_i) \rightarrow$ sum of similarities between x_i and samples in same class

Fitness value \rightarrow adding up all subdata values

- ▶ Host side \rightarrow weight of each sample is calculated
- ▶ Device side \rightarrow data is transmitted to and stored in array *weight*
- ▶ Each thread calculates each subdata in array *SimilarityArray*
- ▶ Each sample is divided into **self-class** and **nonself-class** according to similarity
- ▶ Each thread calculates value of S_{self} and $S_{nonself}$ in parallel

Implementation : Environment

CPU : Intel i7 4th Gen.

- ▶ CPU Cores → 8
- ▶ CPU Clock → 3.40 GHz
- ▶ L3 Cache → 8 MB


GPU : NVIDIA GeForce GT740

- ▶ CUDA Cores → 384
- ▶ Graphics Clock → 993 MHz
- ▶ Memory → 2048 MB DDR3

Implementation

- ▶ C++ Serial program for Normalization of Data
- ▶ C++ Parallel program for Normalization of Data using OpenMP
- ▶ C++ Parallel program for Normalization of Data using CUDA on NVIDIA's GPU

Implementation : Speedup vs. Cores



- ❖ Average Time Required to CUDA Program for Execution → 1085 microsec.
- ❖ Speedup GPU vs CPU → 7.9253


Result & Analysis

- ▶ Experiment
 - ▶ C / C++ programming environment and Linux Operating System
 - ▶ NVIDIA's high performance Tesla K10 GPU
 - ▶ 8 SMs each consisting 192 SPs
 - ▶ SP's clock frequency → 750 MHz
 - ▶ In each block → 65,536 registers
 - ▶ Each Multiprocessor → Thousands of CUDA cores
- ▶ 10-Fold Cross-Validation
 - ▶ Dataset divided into 10 subsets
 - ▶ One is used for testing
 - ▶ Remaining are used for training
 - ▶ Then average performance reported

Result & Analysis

Dataset	Size	NNP _{Seconds}	PNNP _{Seconds}	Speedup Ratio
IONOSPHERE	351 x 34	1283	23	55.78
UKM	403 x 5	1158	15	77.2
VEHICLE	846 x 18	6230	66	94.39
CMC	1473 x 9	19583	219	89.42
SEGMENT	2310 x 19	71198	783	90.93
SATELLITE	6435 x 36	578905	5178	111.8
MAGIC	19020 x 10	8177290	49773	164.29

Result & Analysis



Result & Analysis

	Traditional (Wang et al., 2017)	SoftMax (Wang et al., 2017)	ECOC (Wang et al., 2017)	FCM (Wang et al., 2017)	PNNP
IONOSPHERE	86.18(\pm 6.42)	N/A	N/A	93.44(\pm 5.43)	93.59(\pm1.63)
PARKINSONS	83.12(\pm 16.57)	N/A	N/A	76.75(\pm 17.39)	85.52(\pm9.07)
HABERMAN	59.42(\pm 8.02)	N/A	N/A	55.93(\pm 9.61)	66.79(\pm5.73)
UKM	92.73(\pm 5.03)	85.22(\pm 14.79)	93.26(\pm 2.98)	96.11(\pm 2.03)	97.64(\pm2.01)
VEHICLE	81.88(\pm 2.14)	83.28(\pm2.19)	80.10(\pm 3.12)	78.54(\pm 4.17)	83.19(\pm 2.24)
CMC	50.84(\pm 4.88)	50.49(\pm 2.89)	49.28(\pm 3.16)	52.83(\pm 2.55)	53.08(\pm2.5)
SEGMENT	96.66(\pm 1.79)	96.55(\pm 1.40)	96.16(\pm 1.37)	95.72(\pm 1.69)	96.92(\pm0.90)
VC	72.50(\pm 5.72)	79.98(\pm 6.31)	73.59(\pm 9.61)	81.21(\pm 7.77)	83.02(\pm6.68)
SEEDS	93.32(\pm 6.01)	92.85(\pm 7.15)	94.26(\pm 7.43)	95.18(\pm 5.08)	96.16(\pm3.06)
WINE	97.85(\pm 3.69)	98.33(\pm 3.77)	97.78(\pm 2.88)	98.88(\pm 2.36)	99.49(\pm1.62)

Conclusion

- ▶ To speedup the training process of NNP, particularly for large dataset, we proposed parallel NNP based on NVIDIA's CUDA framework.
- ▶ At CPU Side → Main optimization algorithm performed in serial.
- ▶ At GPU Side → NNP Neural Network Classifier is evaluated parallelly on multiple blocks and all subtasks are performed parallelly using threads.
- ▶ PNNP not only increases speed but also improves measurements.
- ▶ PNNP yields promising that it is able to solve real world problems.
- ▶ In the future, PNNP can be used in medical image analysis and bioinformatics with referenced to big data.

References

- ▶ Anderson, D., Coupland, S., **Parallelisation of fuzzy inference on a graphics processor unit using the compute unified device architecture.** Recent. Progr. Med. 85 (3), 160-165, 2008.
- ▶ Jang, H., Park, A., Jung, K., **Neural network implementation using cuda and openmp.** In: **Digital Image Computing: Techniques and Applications.** pp. 155-161, 2008.
- ▶ Wang, L., Yang, B., Chen, Y., **Improving particle swarm optimization using multilayer searching strategy.** Inform. Sci. 274 (8), 70-94, 2014.
- ▶ Wang, L., Yang, B., Chen, Y., Abraham, A., Sun, H., Chen, Z., Wang, H., **Improvement of neural network classifier using floating centroids.** Knowl. Inf. Syst. 31 (3), 433-454, 2012.

The background features abstract, overlapping green geometric shapes, primarily triangles and polygons, in various shades of green, creating a modern and dynamic visual effect.

Thank You !